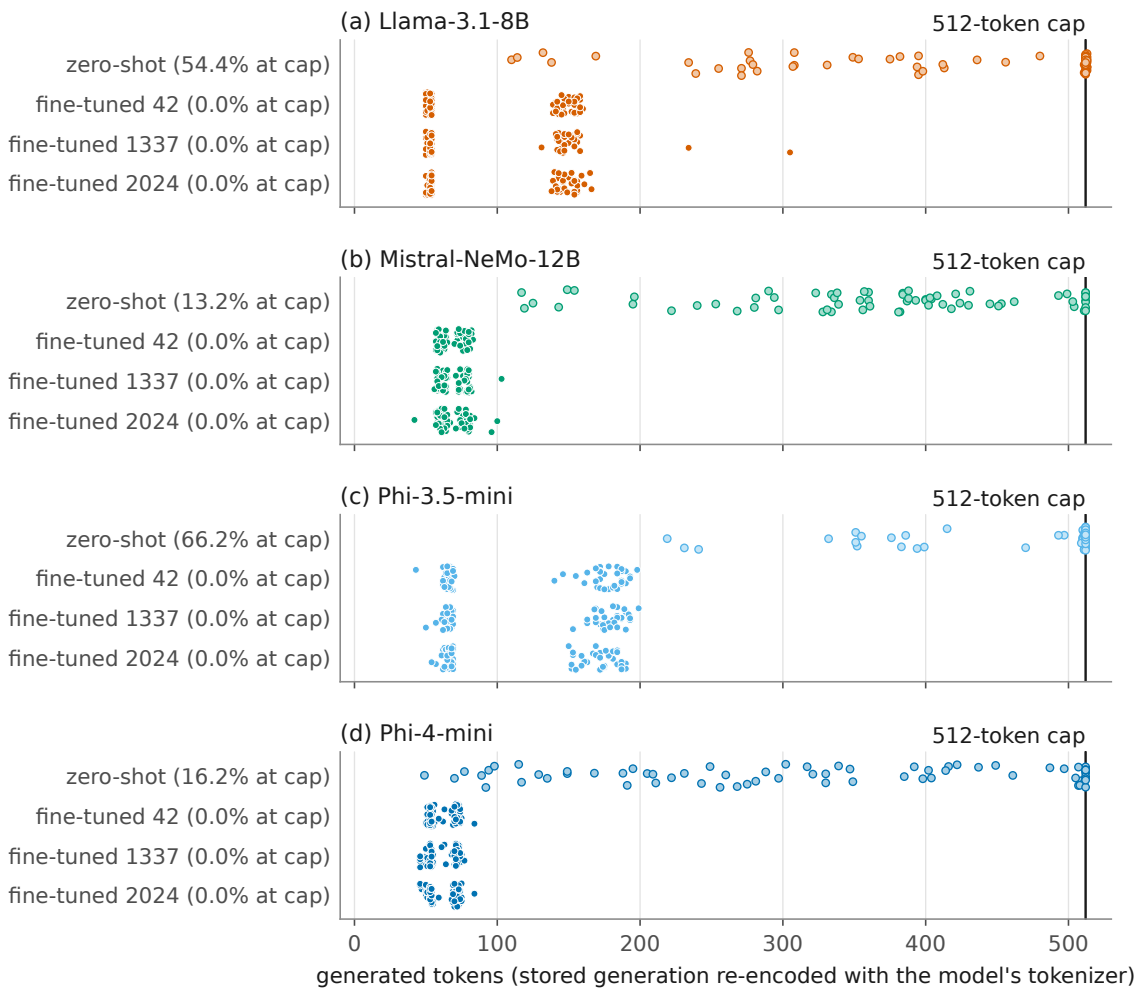
(a) Llama-3.1-8B
512-token cap
zero-shot (54.4% at cap)
fine-tuned 42 (0.0% at cap)
fine-tuned 1337 (0.0% at cap)
fine-tuned 2024 (0.0% at cap)
(b) Mistral-NeMo-12B
512-token cap
zero-shot (13.2% at cap)
fine-tuned 42 (0.0% at cap)
fine-tuned 1337 (0.0% at cap)
fine-tuned 2024 (0.0% at cap)
(c) Phi-3.5-mini
512-token cap
zero-shot (66.2% at cap)
fine-tuned 42 (0.0% at cap)
fine-tuned 1337 (0.0% at cap)
fine-tuned 2024 (0.0% at cap)
(d) Phi-4-mini
512-token cap
zero-shot (16.2% at cap)
fine-tuned 42 (0.0% at cap)
fine-tuned 1337 (0.0% at cap)
fine-tuned 2024 (0.0% at cap)
0
100
200
300
400
500
generated tokens (stored generation re-encoded with the model's tokenizer)